

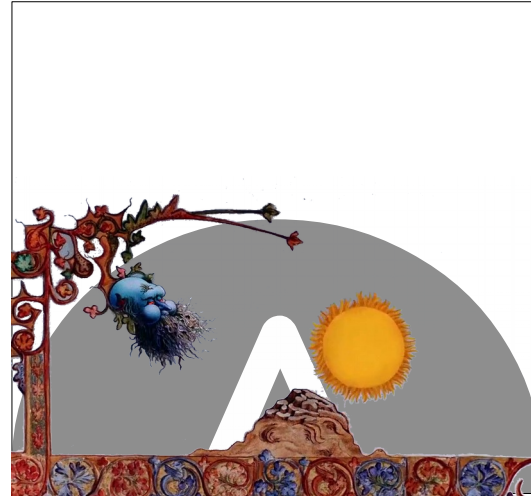
Rehabilitating Pickle

Alex Willmer
Developer, CGI
PyConUK 2018

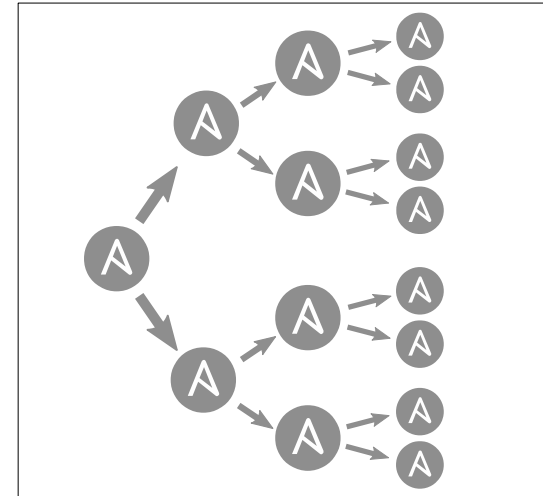
Previously at CGI



DevOps with Ansible
There is much rejoicing
(Yayyy)



Playbooks grow, get slow
Summer becomes Winter



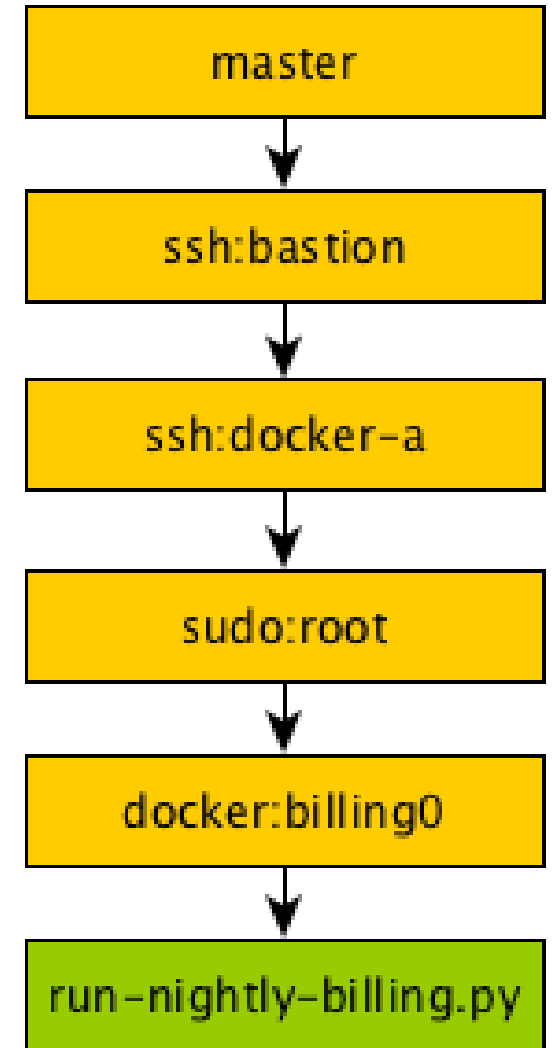
Mitogen for Ansible appears

als
With mitogen my playbook runtime went from 45 minutes to just under 3 minutes. Awesome work!"
The runtime was reduced from 1.5 hours on 4 servers to just under 3 minutes. Thanks!"
h, performance improvement using Mitogen is huge (mentioned before, running with Mitogen enables tasks to take a few seconds). Without Mitogen, the same task takes 19m49s! I'm not even deploying without Mitogen anymore :)"
Works like a charm, thank you for your quick response. He is not kidding about the speed increase. I don't know what kind of dark magic @dmw_83 has. This Mitogen strategy took Clojars' Ansible runs from 45 minutes to 2 minutes. I still can't quite believe it."

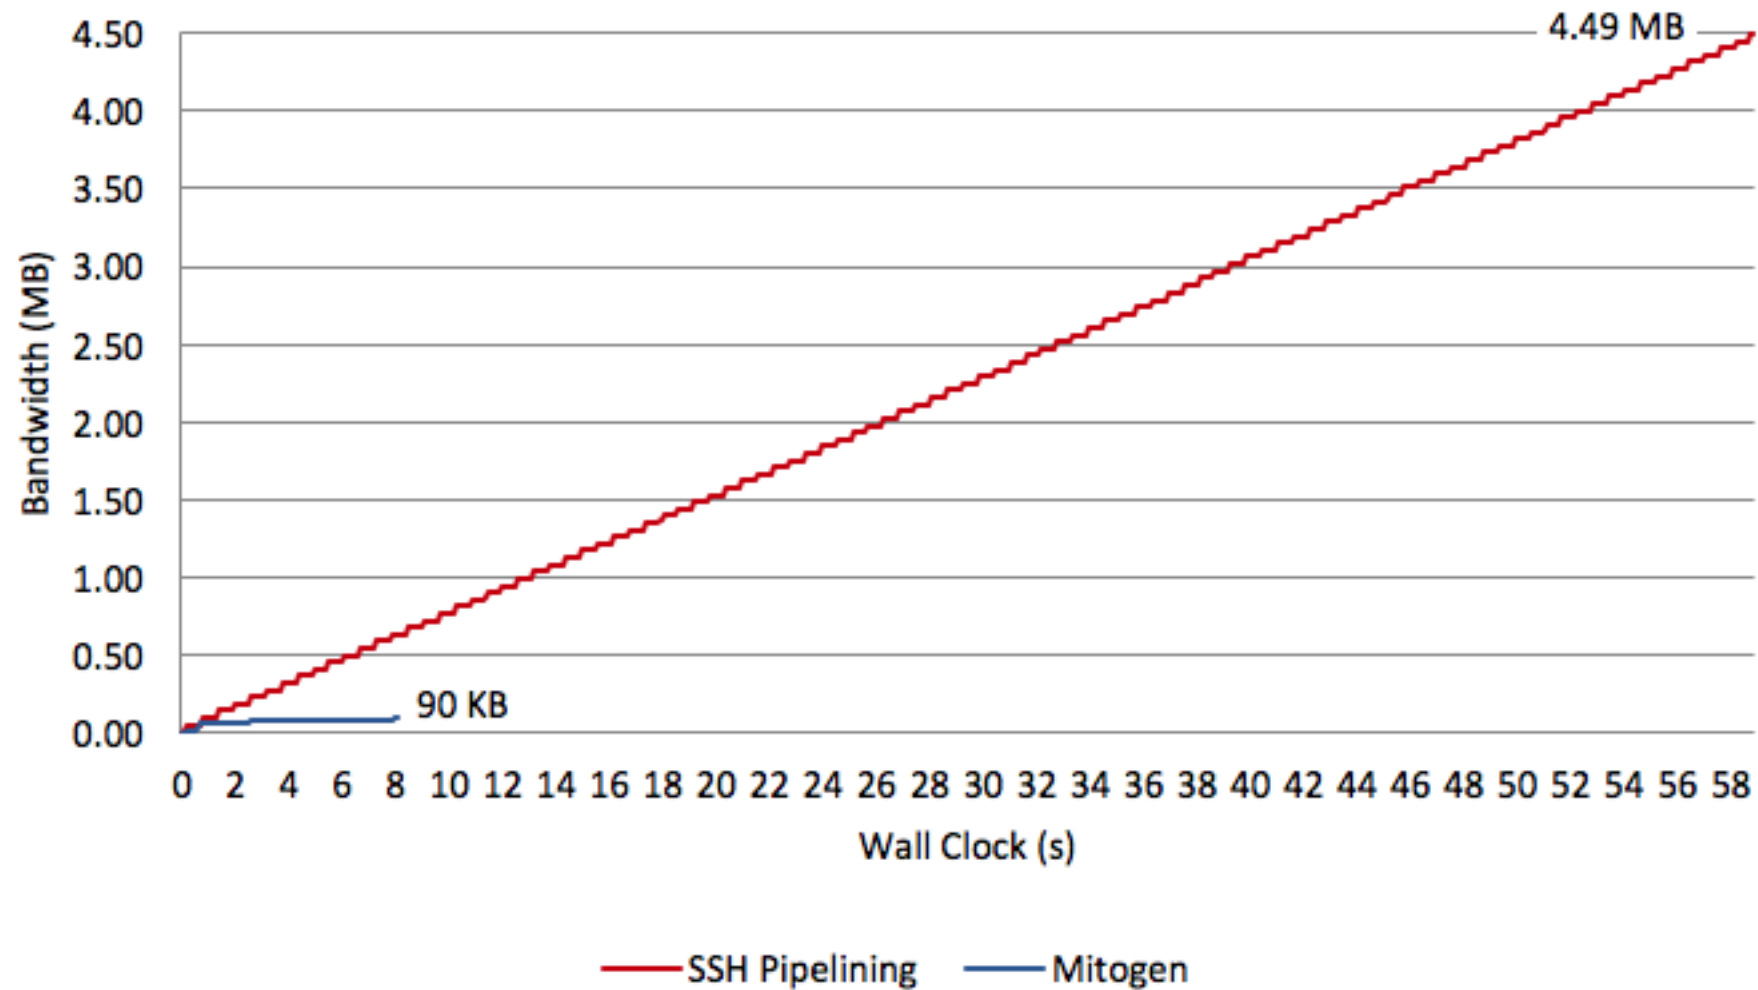
Ansible is fast again ...

Mitogen

```
bastion = router.ssh(  
    hostname='jump-box.mycorp.com')  
  
docker_host = router.ssh(  
    via=bastion_host,  
    hostname='docker-a.prod.mycorp.com')  
  
sudo_account = router.sudo(  
    via=docker_host,  
    username='azure_diamond', password='hunter2')  
  
internal_box = router.docker(  
    via=sudo_account,  
    container='billing0')  
  
internal_box.call(os.system, './run-nightly-billing.py')
```



Mitogen makes Ansible faster



Pickle

```
>>> original = (1.0, 'text', [2,1], {3: None})
```

```
>>> pickle.dumps(original, protocol=0)
```

```
b' (F1.0\nVtext\np0\nn(lp1\nnL2L\nna(dp2\nnL3L\nnNstp3\nn. '
```

```
>>> pickle.loads(_)
```

```
(1.0, 'text', [2,1], {3: None})
```

Why Pickle?

- Preserves nearly any type or class
 - JSON can't do tuple(), set(), complex(), large integers, some dict
- Preserves self referencing data structures
 - JSON doesn't handle recursive structures
- Works everywhere Python does
- Produces compact serializations

**DON'T USE
PICKLE**

Why not Pickle?

Warning: *The pickle module is not secure against erroneous or maliciously constructed data. Never unpickle data received from an untrusted or unauthenticated source.*

[Python documentation » 12.1 pickle — Python object serialization](#)

Why not Pickle?

Warning: *The pickle module is not secure against erroneous or maliciously constructed data. Never unpickle data received from an untrusted or unauthenticated source.*

Python documentation » 12.1 pickle — Python object serialization

CVE-2005-2875	CVE-2005-3008	CVE-2007-1100	CVE-2007-5741	CVE-2008-3143	CVE-2010-4574
CVE-2011-1113	CVE-2011-2520	CVE-2012-4406	CVE-2013-5093	CVE-2013-5942	CVE-2014-0485
CVE-2014-1714	CVE-2014-3539	CVE-2014-5340	CVE-2014-8165	CVE-2015-0692	CVE-2015-0693
CVE-2015-5164	CVE-2015-5242	CVE-2016-3957	CVE-2017-10803	CVE-2018-7889	CVE-2018-14572

– US National Vulnerabilities Database (NVD), vulnerabilities involving "pickle"

`pickle.load()` executes arbitrary code

```
>>> pickle.loads(b"cos\nsystem\n(S'rm -rf'\ntR.")  
0
```

Demo

Restricting globals

```
from pickle import Unpickler, UnpicklingError

class RestrictedUnpickler(Unpickler):
    def find_class(self, module, name):
        if module != "builtins":
            raise UnpicklingError("Denied!")
        elif name not in {"set", "frozenset"}:
            raise UnpicklingError("Still no")
        return getattr(builtins, name)
```

Other attacks

Denial of service

- Protocol downgrade
- Billion Laughs (pickle bombs)
- Unhandled exceptions

Weird machines

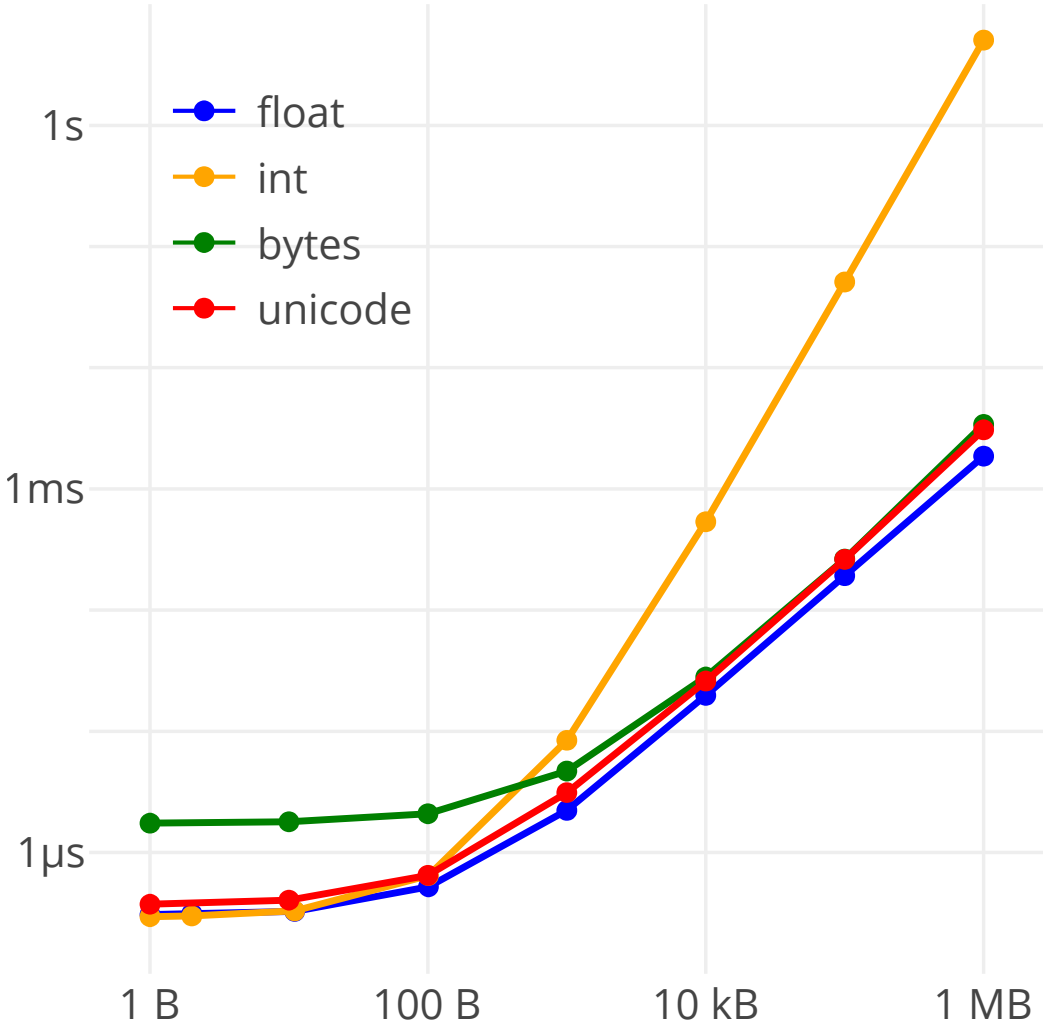
- Unused opcodes (e.g. DUP)
- Parser abuse
- Stack corruption

Data Exposure

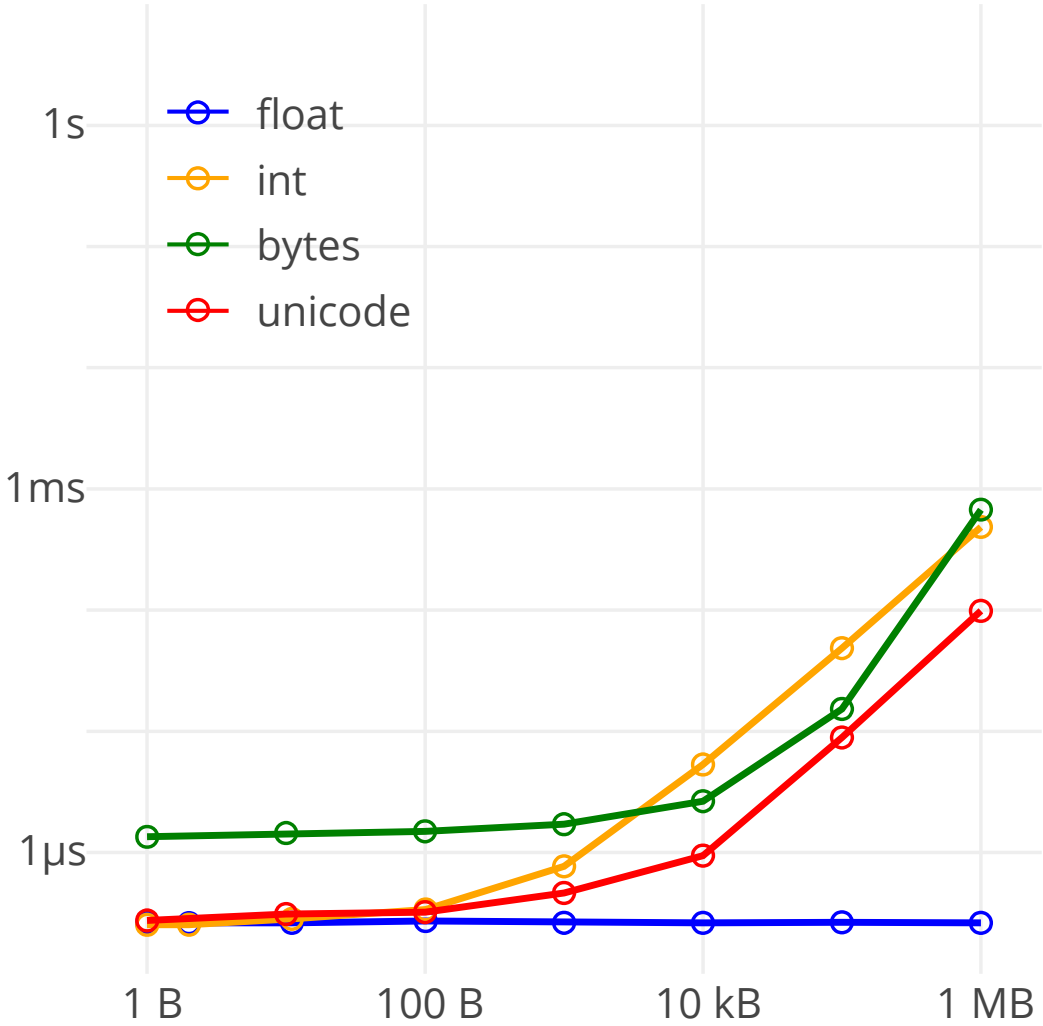
- Object traversal (gadget chains)

Protocol downgrade

Protocol 0



Protocol 1

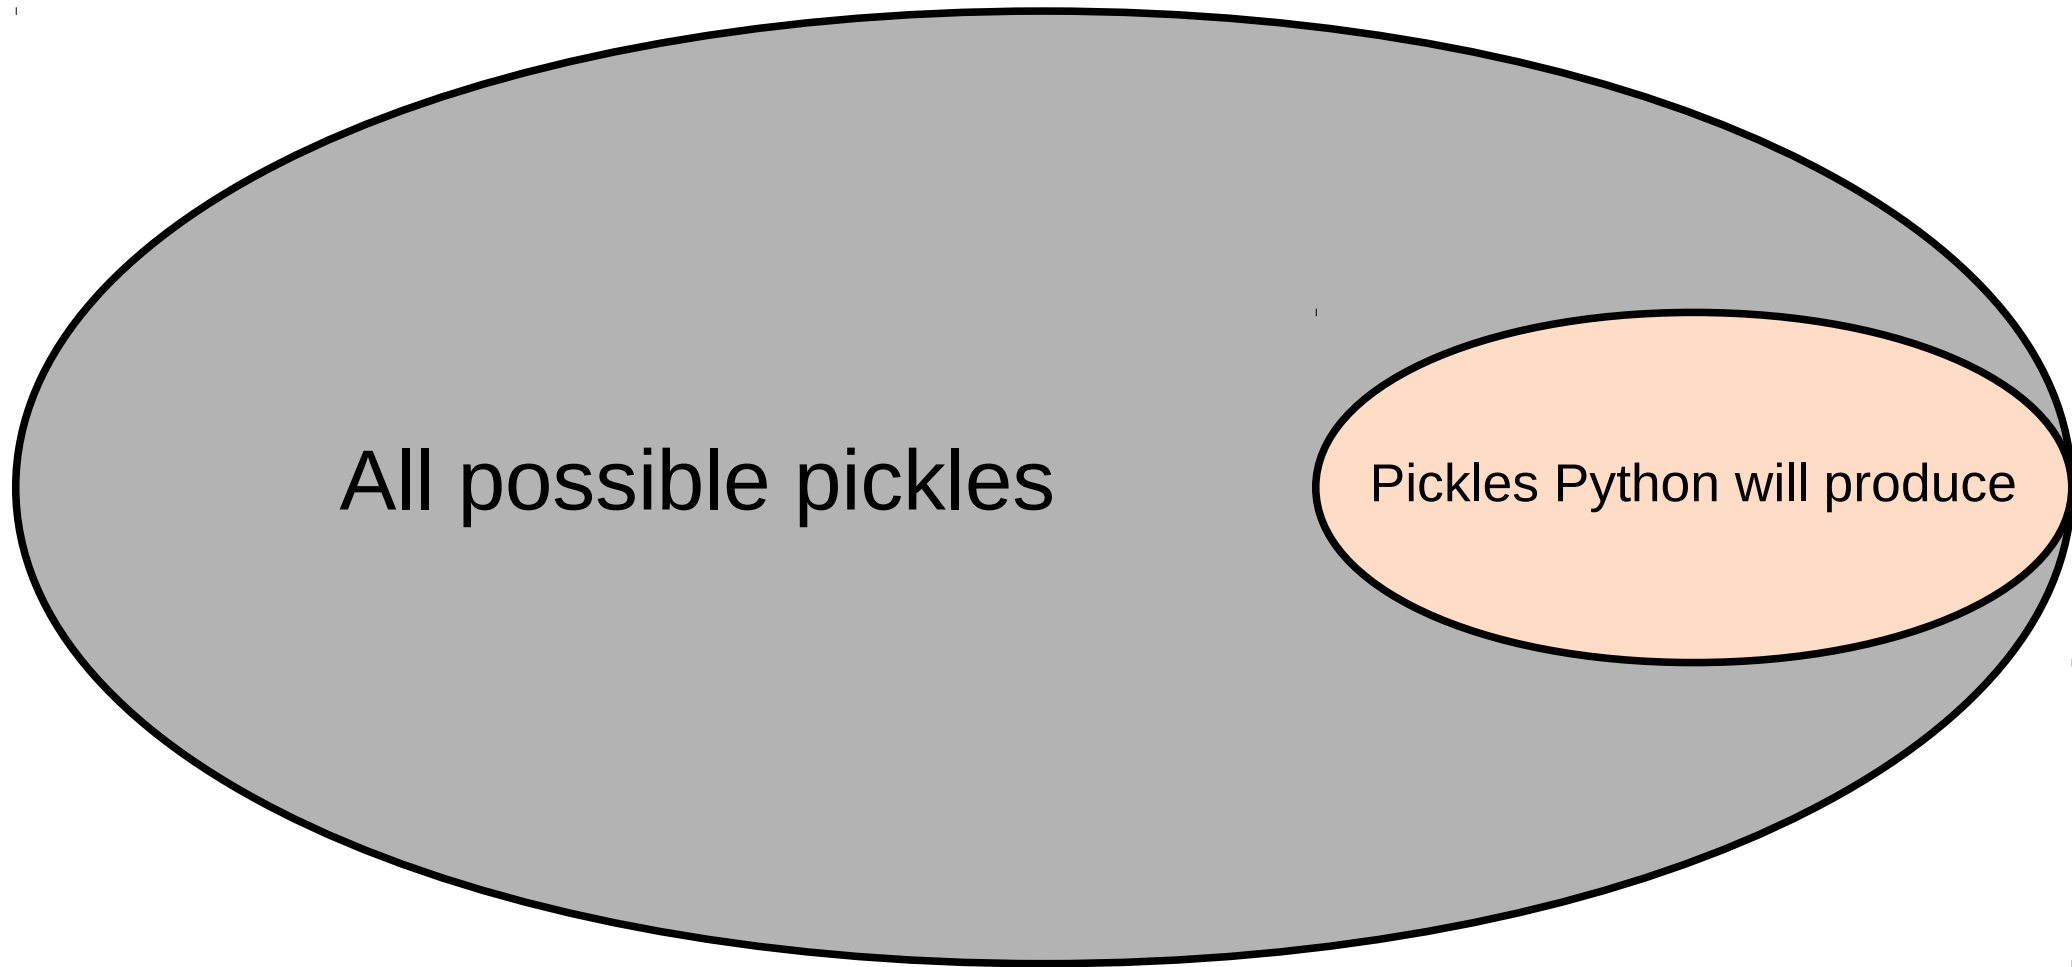


Billion laughs / pickle bombs

```
>>> a = ['lol', 'lol', 'lol', 'lol', 'lol',  
...      'lol', 'lol', 'lol', 'lol', 'lol']  
>>> b = [a,a,a,a,a,a,a,a,a,a]  
>>> c = [b,b,b,b,b,b,b,b,b,b]  
...  
>>> laughs = [i,i,i,i,i,i,i,i,i,i]  
>>> with open('billion-laughs.pk10', 'wb') as f:  
...     pickle.dump(laughs, f)
```

Demo

Weird machines



Other considerations

Don't use pickle for long term storage

“Its automatic, magical behavior shackles you to the internals of your classes in non-obvious ways. You can't even easily tell which classes are baked forever into your pickles. Once a pickle breaks, figuring out why and where and how to fix it is an utter nightmare.”

– Don't use Pickle – use Camel, Eevee

Thank you

Slides, etc.

- github.com/moreati/pickle-fuzz

Mitogen (by David Wilson)

- mitogen.readthedocs.io

Pikara (by Latacora)

- github.com/latacora/pikara

Sour pickles, real world pickle attacks (by Sensepost)

- sensepost.com/blog/2011/blackhat-2011-presentation/
- youtube.com/watch?v=HsZWFMKsM08

Don't use pickle – use camel (by Eevee)

- eev.ee/release/2015/10/15/dont-use-pickle-use-camel/

CGI are hiring



Alex Willmer

Developer, CGI

alex.willmer@cgi.com



@moreati